

Stream processing - part 2

Wojciech Rybak
RTB House

Agenda

- agenda
- intro
- advantages of stream processing
- basic concepts
- design patterns
- fault tolerance
- lambda architecture

What is stream processing?

- “[stream processing is] a type of data processing engine that is designed with infinite data sets in mind” - Tyler Akidau
- infinite == unbounded, always growing
- everything is a stream

Attributes of a data stream

Streams are:

- ordered
- immutable
- replayable

Request-response vs batch vs stream

Request-response:

- low latency
- often blocking
- also known as OLTP - online transaction processing

Batch:

- high latency
- high throughput
- great efficiency
- stale data

Stream:

- nonblocking
- reasonably low latency
- continuous and ongoing

Historical timeline

- 2002: “Models and issues in data stream systems”
- 2003: “TelegraphCQ: continuous dataflow processing”
- ...
- Jan 2011: Kafka
- May 2011: Apache Flink
- Sep 2011: Apache Storm
- ...
- 2013: Amazon Kinesis
- 2015: Google Cloud Pub/Sub + Dataflow
- 2015: Microsoft Azure Stream Analytics

Typical use cases


- complex event processing
- stream analytics
- maintaining materialized views
- search on streams

Real time analysis

- up to date information
- faster decision making
- faster reaction to market fluctuations
- easier detection and addressing of issues
- improved customer experience
- increased agility and optimization


Practical example: bot detector

- certain bots are easy to detect (they make hundreds of actions per minute)
- but difficult to eliminate (they live only for several minutes)
- we can scan a stream of user actions and filter out users with too many actions
- we have actually build it



Short feedback loop when coding

- faster development
- better productivity



Is stream processing always the best approach?

- no

Stream processing concepts

A **stream** is a continuous **body** of **surface water**^[1] flowing within the **bed** and **banks** of a **channel**.

Depending on its location or certain characteristics, a stream may be referred to by a variety of local or regional names. Long large streams are usually called **rivers**, while smaller, less voluminous and more intermittent streams are known as streamlets, brooks or creeks, etc.


The flow of a stream is controlled by three inputs – **surface runoffs** (from **precipitation** or **meltwater**),



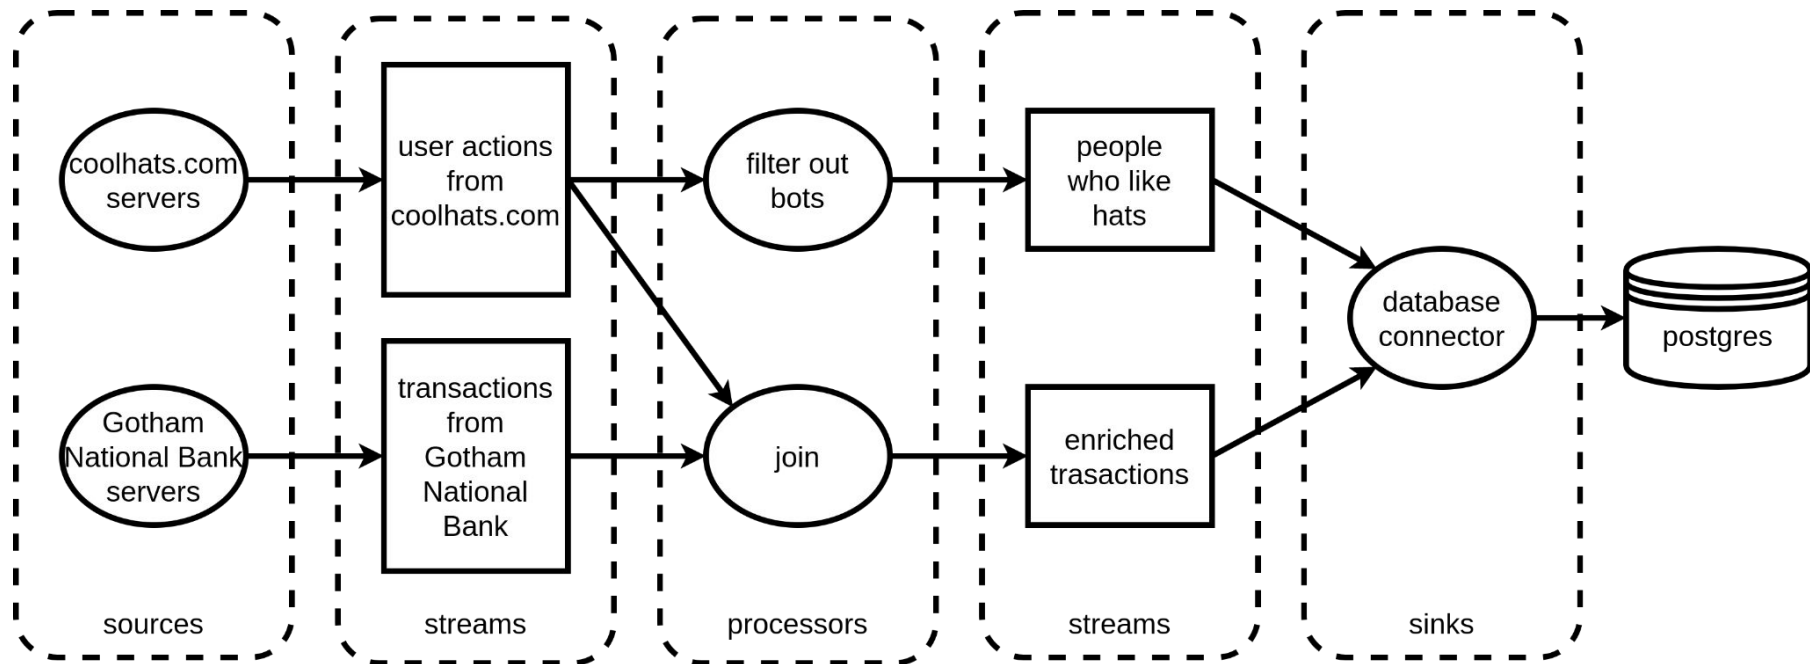
Aubach (Wiehl) in North Rhine-Westphalia, Germany



Topology




Topology




Time

- processing time
- ingestion time
- event time
- sometimes even more (our record: 20 timestamps in a single event)



State

- used to keep track of information about the events
- local: accessible only by a specific instance, within the application
- external: global; outside the application



Stream - table duality

TABLE

id	balance
101	1200
102	-300

UPDATE ... SET balance = 200 WHERE id = 101;
UPDATE ... SET balance = 3000 WHERE id = 101;
UPDATE ... SET balance = 700 WHERE id = 102;
UPDATE ... SET balance = 1200 WHERE id = 101;
UPDATE ... SET balance = -300 WHERE id = 102;

STREAM

101	200
101	3000
102	700
101	1200
102	-300

Stream - table duality


$$state(now) = \int_{t=0}^{now} stream(t) \, dt$$

$$stream(t) = \frac{d \, state(t)}{dt}$$


Processing guarantees

- at most once
- at least once
- exactly once


Design patterns




Single-event processing



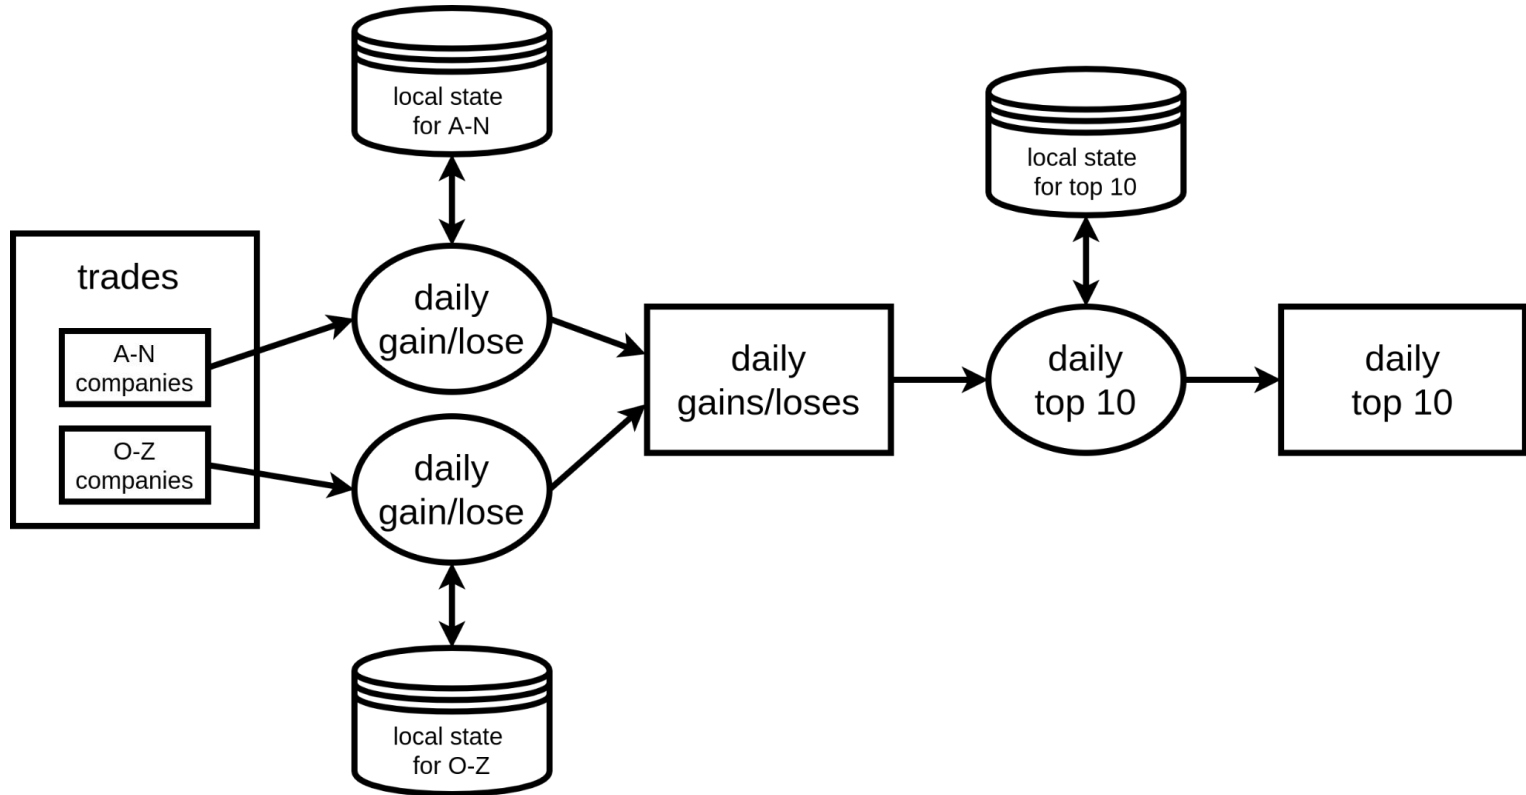
Single-event processing




Local state




Multiphase processing



External lookup (stream - table join)



External lookup (stream - table join)



Streaming join

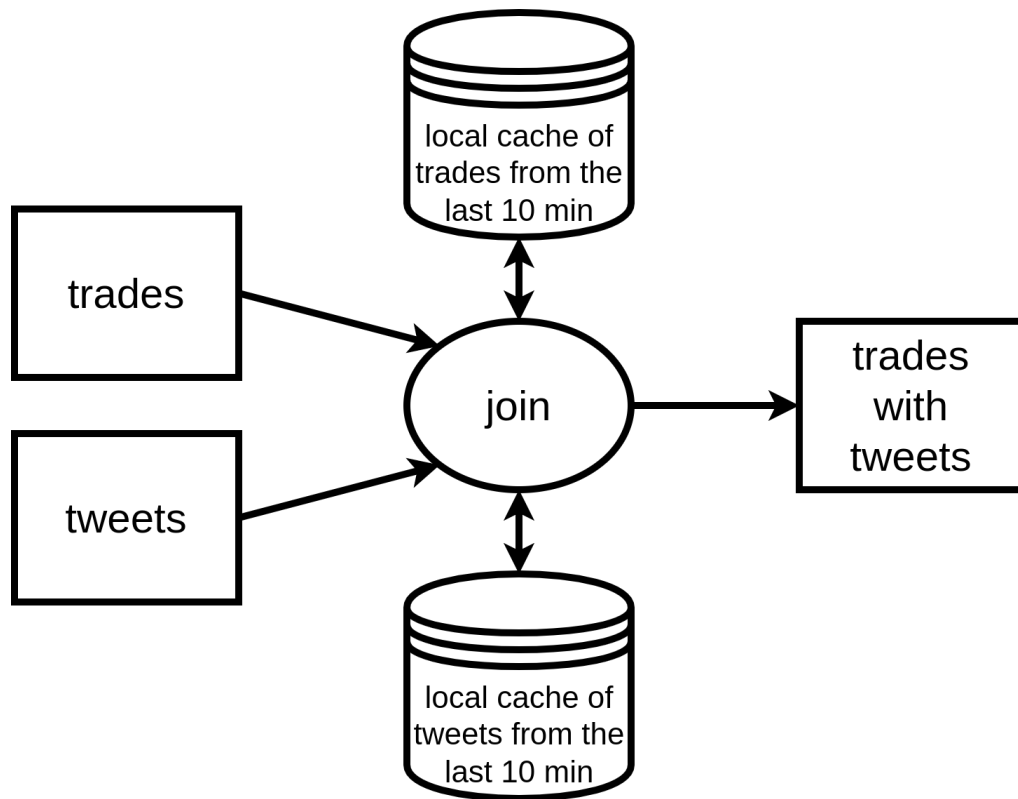


Table - table join

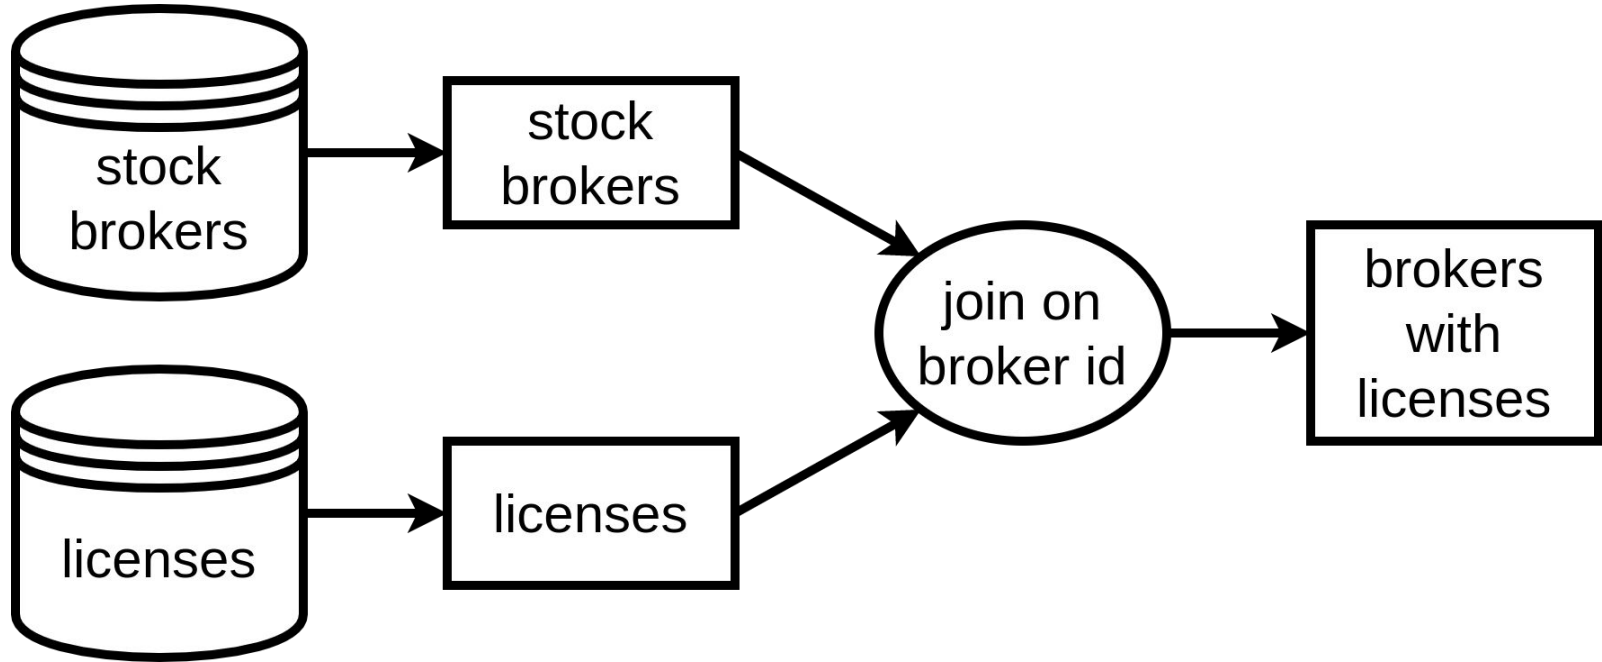



Table - table join




Out of sequence events

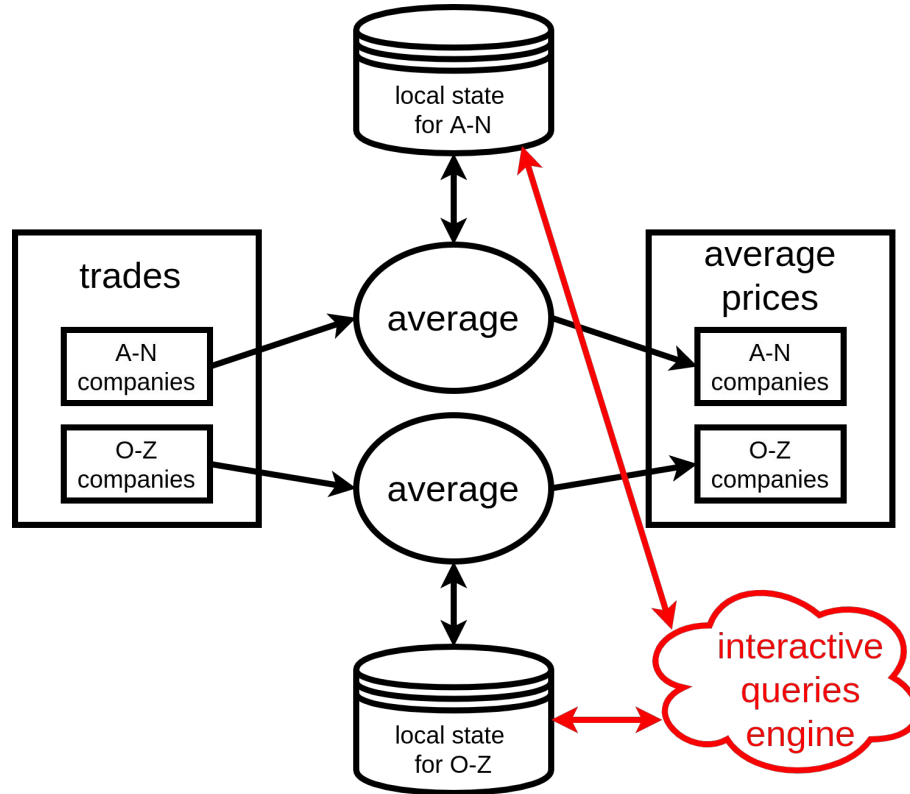
time = 12:01 PM
time = 12:02 PM
time = 12:03 PM
time = 12:04 PM
time = 12:05 PM
time = 12:06 PM
time = 12:07 PM
time = 8:32 AM
time = 8:33 AM
time = 8:34 AM
time = 8:35 AM
time = 12:08 PM
time = 12:09 PM
time = 12:10 PM
time = 12:11 PM
time = 12:12 PM

}
old events
arriving late

Reprocessing




Interactive queries




Time windows




Tumbling (fixed)



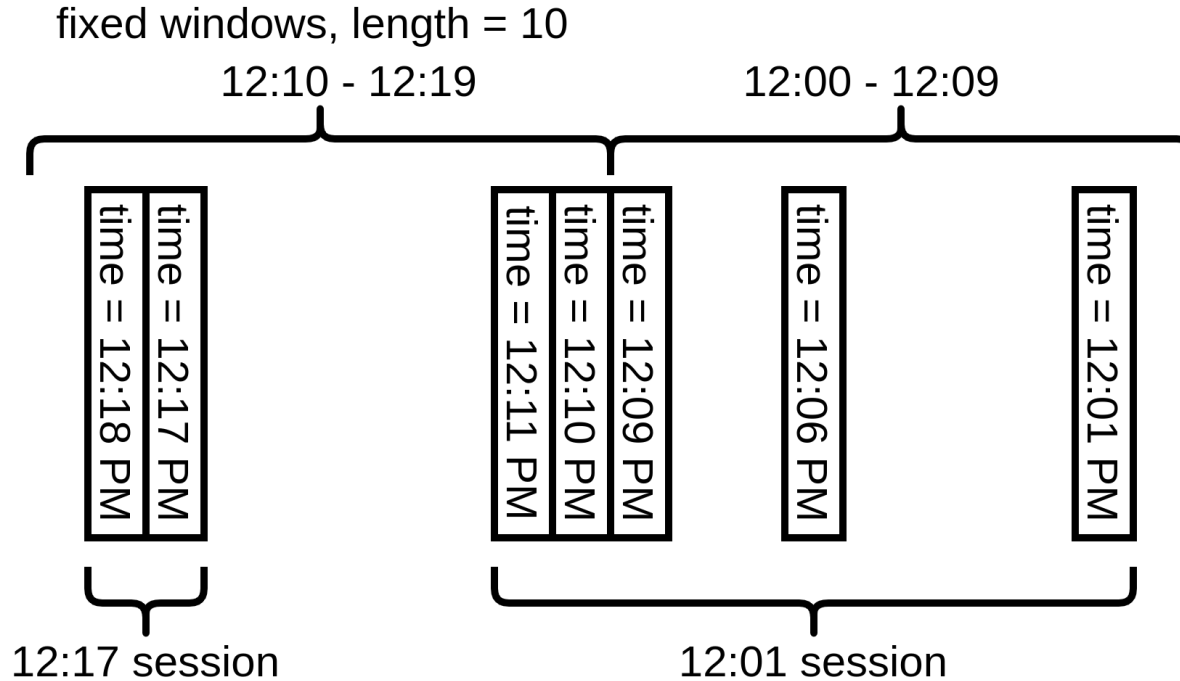
Hopping



Sliding




Session




session windows, max inactivity = 5

Fault tolerance



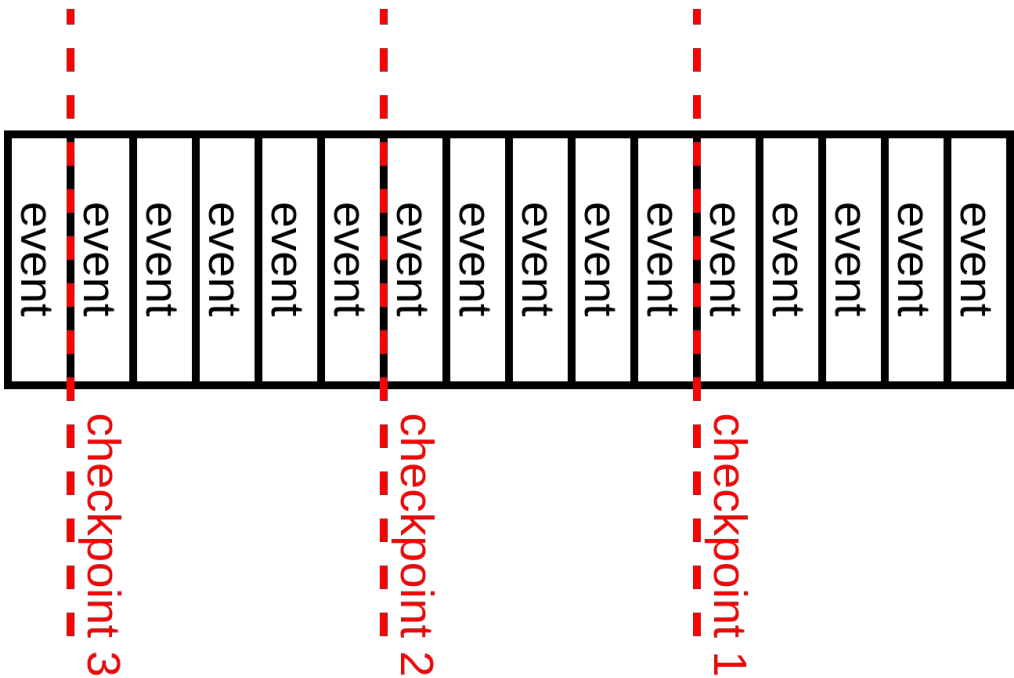
Microbatching

- break the stream into small blocks, and treat each block as a batch process
- used in Spark Streaming



Checkpointing

- periodically generate a checkpoint of state and write it to persistent storage



Atomic commit

- Kafka transactions are of limited use
- X/Open XA (eXtended architecture): global transactions across heterogeneous components, based on two-phase commit protocol
- some large systems support transactions within their environment (e.g. Google Cloud Dataflow)

Idempotence

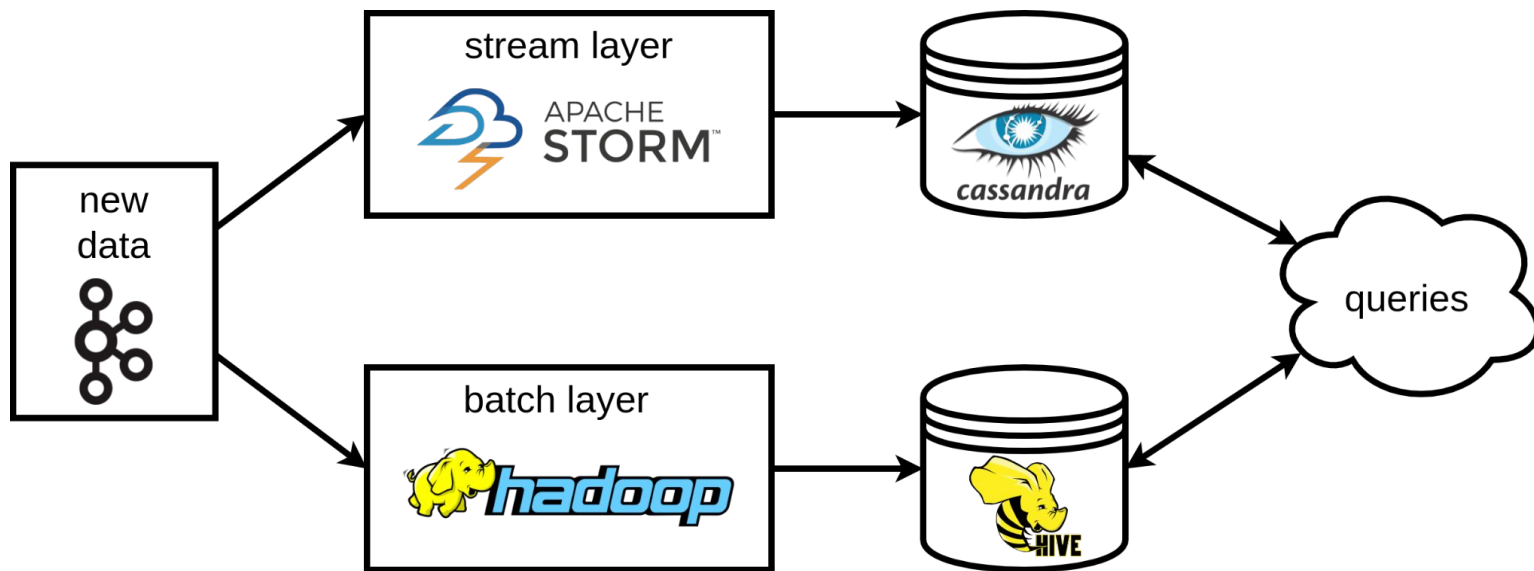
- idempotent operation performed multiple times has the same effect as performed once
- example: “UPDATE table SET processed = true WHERE event_id = 101”
- at least once + idempotent operations = effectively once

Rebuilding state

- applications must be able to restore the state after a failure
- option 1: keep the state in a remote database
- option 2: keep the state local, and replicate it periodically
- option 3: do nothing, and after a failure simply replay the input stream from start and rebuild the state

Lambda architecture

- “How to beat the CAP theorem” - Nathan Marz, 2011



Benefits & challenges


Positives:

- best of two worlds: low latency and possibility to process historical data
- algorithmic flexibility
- easy ad-hoc analysis

Negatives:

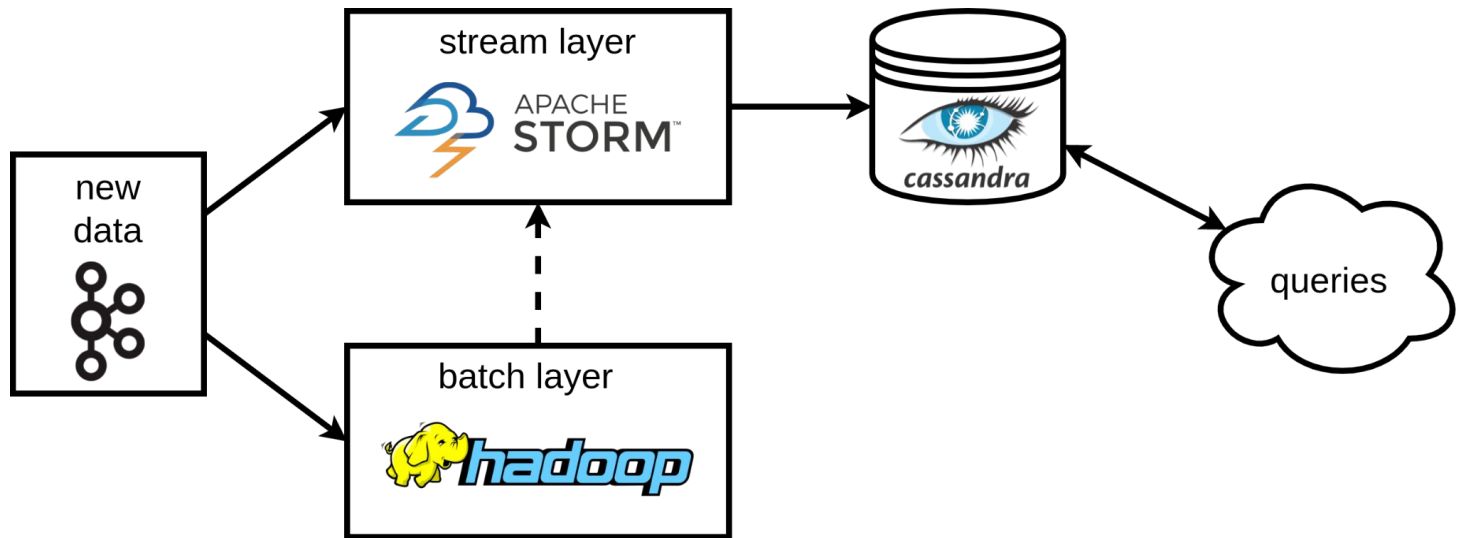
- inherent complexity
- maintaining two codebases
- synchronization between the layers is needed

Practical example: stats counter



Kappa architecture

- “Questioning the Lambda Architecture” - Jay Kreps, 2014



Summary

Two things to remember:

- strengths and weaknesses of data streams
- basic concepts and techniques

